

```
In [32]: import numpy as np
import pandas as pd
df = pd.read_csv("Churn_Modelling.csv")
df
```

```
Out[32]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	10134
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	11254
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	11393
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	9382
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	7908
...
9995	9996	15606229	Obijaku	771	France	Male	39	5	0.00	2	1	0	9627
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	1	1	1	10169
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	0	1	4208
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	2	1	0	9288
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	1	1	0	3819

10000 rows × 14 columns

3. visualizations

```
In [2]: import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

i)Univariate Analysis


```
In [3]: df[['CustomerId', 'Surname', 'CreditScore', 'Geography', 'Age', 'Tenure']].describe()
```

```
Out[3]:
```

	CustomerId	CreditScore	Age	Tenure
count	1.000000e+04	10000.000000	10000.000000	10000.000000
mean	1.569094e+07	650.528800	38.921800	5.012800
std	7.193619e+04	96.653299	10.487806	2.892174
min	1.556570e+07	350.000000	18.000000	0.000000
25%	1.562853e+07	584.000000	32.000000	3.000000
50%	1.569074e+07	652.000000	37.000000	5.000000
75%	1.575323e+07	718.000000	44.000000	7.000000
max	1.581569e+07	850.000000	92.000000	10.000000

```
In [4]: sns.histplot(df.Age, kde=True)
```


```
Out[4]: <AxesSubplot:xlabel='Age', ylabel='Count'>
```



```
In [5]: # plot count plot for the gender column
sns.countplot(df.Gender)
```

C:\Users\User\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()

```
Out[5]: <AxesSubplot:xlabel='Gender', ylabel='count'>
```



ii) Bivariate Analysis

```
In [6]: df[['CustomerId', 'Surname', 'CreditScore', 'Geography', 'Gender', 'Age']].corr()
```


```
Out[6]:
```

	CustomerId	CreditScore	Age
CustomerId	1.000000	0.005308	0.009497
CreditScore	0.005308	1.000000	-0.003965
Age	0.009497	-0.003965	1.000000

```
In [9]: sns.scatterplot(df.CreditScore, df.Age)
plt.ylim(0, 100)
```

C:\Users\User\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()


```
Out[9]: (0.0, 100.0)
```



iii)Multivariate Analysis

```
In [8]: sns.pairplot(data =df[['CustomerId','Geography','Gender','CreditScore','Age','Balance']],hue = 'Balance')
```

```
Out[8]: <seaborn.axisgrid.PairGrid at 0x29ee2462310>
```



4. Descriptive Statistics

```
In [9]: #mode
df['Age'].mode()
```

```
Out[9]: 0    37
        Name: Age, dtype: int64
```

```
In [10]: #calculation of the mean (for Age)
df["Age"].mean()
```

```
Out[10]: 38.9218
```

```
In [11]: #calculation of the mean and round the result(for Age)
round(df["Age"].mean(), 2)
```

```
Out[11]: 38.92
```

```
In [12]: #calculation of the median(for Age)
df["Age"].median()
```

```
Out[12]: 37.0
```

```
In [33]: df.columns
```

```
Out[33]: Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
                'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
                'IsActiveMember', 'EstimatedSalary', 'Exited'],
                dtype='object')
```

```
In [14]: df["NumOfProducts"].value_counts()
```

```
Out[14]: 1    5084
         2    4590
         3     266
         4      60
         Name: NumOfProducts, dtype: int64
```

```
In [34]: df.dtypes
```

```
Out[34]: RowNumber      int64
         CustomerId     int64
         Surname        object
         CreditScore     int64
         Geography      object
         Gender         object
         Age            int64
         Tenure         int64
         Balance        float64
         NumOfProducts  int64
         HasCrCard      int64
         IsActiveMember int64
         EstimatedSalary float64
         Exited         int64
         dtype: object
```

```
In [35]: df.head()
```

```
Out[35]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	113931.57
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93826.63
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10

```
In [17]: df.describe()
```

```
Out[17]:
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.00000	10000.000000	10000.000000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	1.530200	0.70550	0.515100	100090.239881
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	0.581654	0.45584	0.499797	57510.492818
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.00000	0.000000	11.580000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.00000	0.000000	51002.110000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.000000	1.00000	1.000000	100193.915000
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	2.000000	1.00000	1.000000	149388.247500
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	4.000000	1.00000	1.000000	199992.480000

5. Handling Missing values

```
In [18]: df.isna().any()
```

```
Out[18]: RowNumber      False
         CustomerId     False
         Surname        False
         CreditScore     False
         Geography      False
         Gender         False
         Age            False
         Tenure         False
         Balance        False
         NumOfProducts  False
         HasCrCard      False
         IsActiveMember False
         EstimatedSalary False
         Exited         False
         dtype: bool
```

```
In [19]: df.isnull().sum()
```

```
Out[19]: RowNumber      0
CustomerId      0
Surname         0
CreditScore     0
Geography       0
Gender          0
Age             0
Tenure          0
Balance         0
NumOfProducts  0
HasCrCard       0
IsActiveMember  0
EstimatedSalary 0
Exited          0
dtype: int64
```

```
In [21]: df.notnull()
```

```
Out[21]:
```


	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	True	True	True	True	True	True	True	True	True	True	True	True	True
1	True	True	True	True	True	True	True	True	True	True	True	True	True
2	True	True	True	True	True	True	True	True	True	True	True	True	True
3	True	True	True	True	True	True	True	True	True	True	True	True	True
4	True	True	True	True	True	True	True	True	True	True	True	True	True
...
9995	True	True	True	True	True	True	True	True	True	True	True	True	True
9996	True	True	True	True	True	True	True	True	True	True	True	True	True
9997	True	True	True	True	True	True	True	True	True	True	True	True	True
9998	True	True	True	True	True	True	True	True	True	True	True	True	True
9999	True	True	True	True	True	True	True	True	True	True	True	True	True

10000 rows × 14 columns

6. Finding and replacing the outliers


```
In [22]: import seaborn as sns
sns.boxplot(x=df['Age'])
```

```
Out[22]: <AxesSubplot:xlabel='Age'>
```



```
In [23]: sns.boxplot(x=df['Tenure'])
```

```
Out[23]: <AxesSubplot:xlabel='Tenure'>
```



7. Check for categorical columns and perform encoding

```
In [24]: import pandas as pd
df = pd.read_csv("Churn_Modelling.csv", header=None)
```

```
In [25]: cols = df.columns
num_cols = df._get_numeric_data().columns
```

```
In [26]: num_cols
```

```
Out[26]: Int64Index([], dtype='int64')
```

```
In [27]: list(set(cols) - set(num_cols))
```

```
Out[27]: [0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13]
```

8. Split the data into dependent and independent variables

```
In [36]: x = df.drop('Exited', axis=1)
y = df['Exited']
```

```
In [37]: x.head()
```

```
Out[37]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	113931.57
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93826.63
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10

```
In [38]: y.head()
```

```
Out[38]: 0    1
1    0
2    1
3    0
4    0
Name: Exited, dtype: int64
```

9. Scale the independent variables

```
In [39]: from sklearn import linear_model
from sklearn.preprocessing import StandardScaler
scale = StandardScaler()
```

```
In [40]: X = df[['Balance', 'Tenure']]

scaledX = scale.fit_transform(X)
print(scaledX)

[[-1.22584767 -1.04175968]
 [ 0.11735002 -1.38753759]
 [ 1.33305335  1.03290776]
 ...
 [-1.22584767  0.68712986]
 [-0.02260751 -0.69598177]
 [ 0.85996499 -0.35020386]]
```

10. Split the data into training and testing

```
In [41]: from sklearn.model_selection import train_test_split
```

```
In [42]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,random_state=0)
```

```
In [43]: print('X Train shape:{},Y.Train SHape:{}'.format(x_train.shape,y_train.shape))

X Train shape:(8000, 13),Y.Train SHape:(8000,)
```

```
In [44]: print('X Test Shape :{},Y Test SHape:{}'.format(x_test.shape,y_test.shape))

X Test Shape :(2000, 13),Y Test SHape:(2000,)
```